

CS 61C Fall 2018

RISC-V Addressing and Caches

Discussion 5: September 24, 2018

1 RISC-V Addressing

We have several *addressing modes* to access memory (immediate not listed):

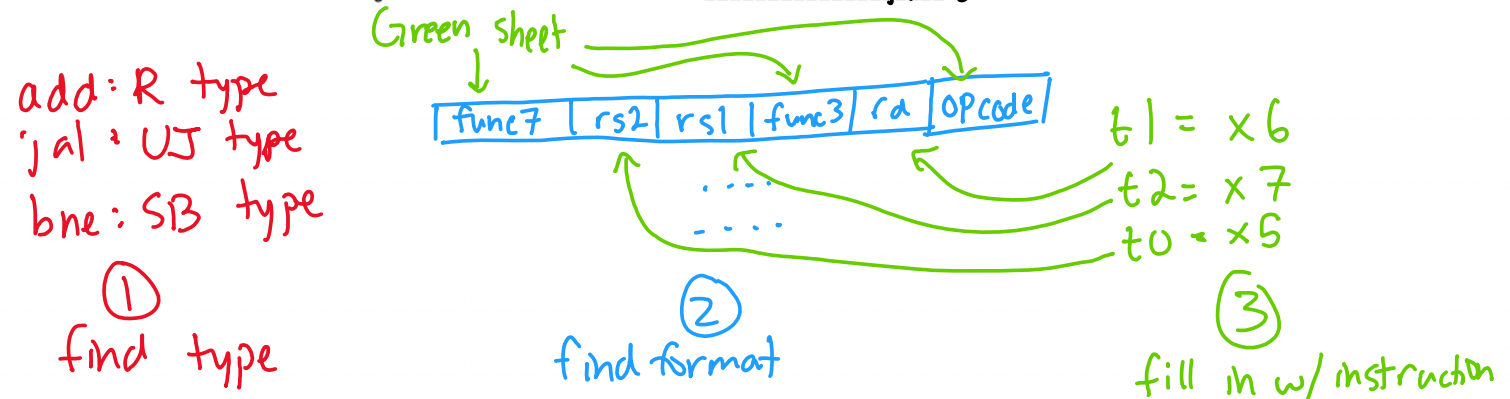
1. Base displacement addressing adds an immediate to a register value to create a memory address (used for lw, lb, sw, sb).
2. PC-relative addressing uses the PC and adds the immediate value of the instruction (multiplied by 2) to create an address (used by branch and jump instructions).
3. Register Addressing uses the value in a register as a memory address (jr)

★ 1.1 What is range of 32-bit instructions that can be reached from the current PC using a branch instruction? 12 bits $\Rightarrow [-2^{11}, 2^{11}-1]$ half words
 $\Rightarrow [-2^{10}, 2^{10}-1]$ words

★ 1.2 What is the range of 32-bit instructions that can be reached from the current PC using a jump instruction? 20 bits $\Rightarrow [-2^{19}, 2^{19}-1]$ half words
 $\Rightarrow [-2^{18}, 2^{18}-1]$ words

★ 1.3 Given the following RISC-V code (and instruction addresses), fill in the blank fields for the following instructions (you'll need your RISC-V green card!).

1	0x002cfff0: loop: add t1, t2, t0	0	5	7	0	6	0x33
2	0x002cfff4: jal ra, foo	0	0x14	0	0	1	0x6F
3	0x002cfff8: bne t1, zero, loop	1	0x3F	0	6	1	0xC63
4	...						
5	0x002cfff2c: foo: jr ra						



(3½) find offsets

jal: 0x002cfff2c
 - 0x002cfff04
 = 0x28 = 0b 0010 1000 = 20 bits

half words

20 bits

half words

bit #

2 Understanding T/I/O

When working with caches, we have to be able to break down the memory addresses we work with to understand where they fit into our caches. There are three fields:

Tag - Used to distinguish different blocks that use the same index - Number of bits: leftovers *leftover bits*

Index - The set that this piece of memory will be placed in - Number of bits: $\log_2(\# \text{ of indices})$

Offset - The location of the byte in the block - Number of bits: $\log_2(\text{size of block})$

- 2.1 Assume we have a direct-mapped byte-addressed cache with capacity 32B and block size of 8B. Of the 32 bits in each address, which bits do we use to find the index of the cache to use?

$$\text{Offset} = \log 2^3 = 3 \quad \frac{32\text{B}}{8\text{B}} = 4 \text{ blocks} \Rightarrow \text{index} = \log 2^2 = 2$$

- 2.2 Which bits are our tag bits? What about our offset?

→ Tag: 31:5 Index: 4:3 Offset: 2:0



- 2.3 Classify each of the following byte memory accesses as a cache hit (H), cache miss (M), or cache miss with replacement (R). It is probably best to try drawing out the cache before going through so that you can have an easier time seeing the replacements in the cache. The following white space is to do this:

Address	T/I/O	Hit, Miss, Replace
0x00000004 ... 0100	0 0 4	M, Comp
0x00000005 ... 0101	0 0 5	H
0x00000068 0110 1000	3 1 0	M, Comp
0x000000C8 1100 1000	6 1 0	R, Comp
0x00000068 0110 1000	3 1 0	R, conflict
0x000000DD 1101 1101	6 3 5	M, Comp
0x00000045 0100 0101	2 0 5	R, Comp
0x00000004 ... 0100	0 0 4	R, Conflict
0x000000C8 1100 1000	6 1 0	R, Conflict

Cache

Index	Tag	Contents	valid
00	0100	---	1
01	0101	---	1
10			0
11	6	---	1

8B of data here

gray area between Conflict & capacity miss
(it's both full & a conflict)

Offset: find data in block
index: find block in cache
tag: ensure we're looking @ correct block

Cache size
block size

$\log(\# \text{ indices})$
 $\log(\text{Block size})$
for direct mapped cache

31 ... 5 43 2 1 0
T I O

use TIO from before

3 The 3 C's of Misses

3.1 Classify each M and R above as one of the 3 types of misses described below:

- I. Compulsory: First time you ask the cache for a certain block. A miss that must occur when you first bring in a block. Reduce compulsory misses by having a longer cache lines (bigger blocks), which bring in the surrounding addresses along with our requested data. Can also pre-fetch blocks beforehand using a hardware prefetcher (a special circuit that tries to guess the next few blocks that you will want).
- II. Conflict: Occurs if you hypothetically went through the ENTIRE string of accesses with a fully associative cache and wouldn't have missed for that specific access. Increasing the associativity or improving the replacement policy would remove the miss.
- III. Capacity: The only way to remove the miss is to increase the cache capacity, as even with a fully associative cache, we had to kick a block out at some point.

TL ; DR

first time
Seen

mostly direct mapped

can organize better

cache full, need to enlarge

Note: There are many different ways of fixing misses. The name of the miss doesn't necessarily tell us the best way to reduce the number of misses.

4 Extra Practise

In the following diagrams, each blank box represents 1 byte (8 bits) of data. All of memory is byte addressed. Let's say we have a 8192KiB cache with an 128B block size, how many bits are in tag, index, and offset? What parts of the address of 0xFEEDF00D fit into which sections?

$2^{13} \cdot 2^{10} / 2^7$ $\log 2^7 = 7$

$32 - 16 - 7 = 9$

	Tag	Index	Offset
Number of bits	9	16	7
Bits of address	11 11 1110 1	110 1101 111 100000	000 1101

0x 1111 1110 1110 1101 111 0000 0000 1101

T I O

4.2 Now fill in the table below. Assume that we have a write-through cache, so the number of bits per row includes only the cache data, the tag, and the valid bit.

	Address size (bits)	Cache Size	Block Size	Tag Bits	Index Bits	Offset Bits	Bits per row
1.	16	4KiB	4B	4	10	2	37
2.	32	32KiB	16B	17	11	4	146
3.	32	64KiB	16B	16	12	4	145
4.	64	2048KiB	128B	43	14	7	1068

Data + tag + valid

only 1 less tag bit

1. O: $\log 4 = 2$ I: $\log \frac{2^{12}}{2^2} = 10$ T: $16 - 10 - 2 = 4$ $32 + 4 + 1 = 37$
2. O: $\log 16 = 4$ I: $\log \frac{2^{15}}{2^4} = 11$ T: $32 - 11 - 4 = 17$ $128 + 17 + 1 = 146$
3. O: $32 - 16 - 12 = 4$ Block: $2^4 = 16B$ size: $2^{12} \cdot 16B = 2^{16} = 64KiB$
4. Block: $\frac{2^{11} \cdot 2^{10}}{2^{14}} = 2^7 = 128B$ O: $\log 2^7 = 7$ T: $64 - 14 - 7 = 43$